

Project #1: Malloc Library

Part 1

ECE 650 – Spring 2025

Due on 1/25 11:59 pm

General Instructions

1. You will work individually on this project.
2. The code for this assignment should be developed and tested in a UNIX-based environment, specifically the Duke Linux environment available at login.oit.duke.edu or <https://vcm.duke.edu/>.
3. You must follow this assignment spec carefully, and turn in everything that is asked (and in the proper formats, as described). Due to the large class size, this is required to make grading more efficient.
4. You should plan to start early on this project and make steady progress over time. It will take time and careful thought to work through the assignment.
5. You have a chance to preview your program results tested by our auto-grader if you submit your homework on Gradescope 2 days before the due date. The autograder will release your score 48 hours before the deadline. The pre-release is just a feedback for you and isn't meant to help you debug, so please do not ask your TA for what test cases you have failed.

Introduction

For this assignment, you will implement your own version of several memory allocation functions from the C standard library (actually you will have the chance to implement and study several different versions as described below). Your implementation is to be done in C code.

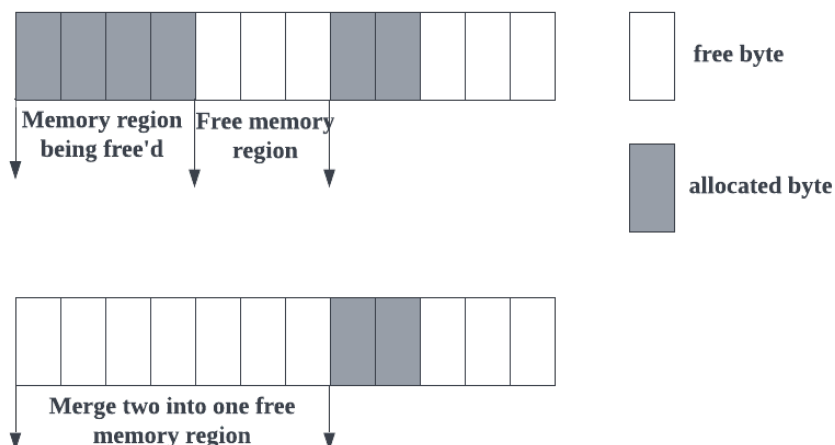
The C standard library includes 4 malloc-related library functions: `malloc()`, `free()`, `calloc()`, and `realloc()`. In this assignment, you only need to implement versions of `malloc()` and `free()`:

```
void *malloc(size_t size);  
void free(void *ptr);
```

Please refer to the man pages for full descriptions of the expected operation for these functions. Essentially, `malloc()` takes in a size (number of bytes) for a memory allocation, locates an address in the program's data region where there is enough space to fit the specified number of bytes, and returns this address for use by the calling program. The `free()` function takes an address (that was returned by a previous `malloc` operation) and marks that data region as available again for use.

The submission instructions at the end of this assignment description provide specific details about what code files to create, what to name your new versions of the malloc functions, etc.

As you work through implementing `malloc()` and `free()`, you will discover that as memory allocations and deallocations happen, you will sometimes free a region of memory that is adjacent to other also free memory region(s). **Your implementation is *required* to coalesce in this situation by merging the adjacent free regions into a single free region of memory.** Similarly, it is also required to split free regions if the ideal free region is larger than requested size. The following graph illustrates the coalesce of the regions.




In this assignment, you will develop a malloc implementation and study different allocation policies. In Homework 2, you will make this implementation thread-safe.

Study of Memory Allocation Policies

Your task is to implement 2 versions of malloc and free, each based on a different strategy for determining the memory region to allocate. The two strategies are:

1. **First Fit:** Examine the free space tracker (e.g. free list), and allocate an address from the first free region with enough space to fit the requested allocation size.
2. **Best Fit:** Examine all of the free space information, and allocate an address from the free region which has the smallest number of bytes greater than or equal to the requested allocation size.

The following picture illustrates how each strategy would operate (assuming free regions are traversed in a left to right order) for a malloc() request of 2 bytes:



Requirements

Malloc implementations

To implement your allocation strategies, you will create 4 functions:

```
//First Fit malloc/free
void *ff_malloc(size_t size);
void ff_free(void *ptr);

//Best Fit malloc/free
void *bf_malloc(size_t size);
void bf_free(void *ptr);
```

Note, that in all cases, a policy to minimize the size of the process's data segment should be used. In other words, only if there is no free space that fits an allocation request, then `sbrk()` should be used to create that space. However, you do not need to perform any type of garbage collection (e.g. reducing the size of the process's data segment, even if allocations at the top of the data segment have been freed).

On `free()`, your implementation is required to merge the newly freed region with any currently free adjacent regions. In other words, your bookkeeping data structure should not contain multiple adjacent free regions, as this would lead to poor region selection during `malloc`. Since we are managing the memory space by ourselves, there is no need to call the original `free()` function.

Performance study report

In addition to implementing these `malloc` functions, you are tasked to conduct a performance study of the `malloc()` performance with different allocation policies. Several programs for experimentation will be provided that perform `malloc()` and `free()` requests with different patterns (e.g. frequencies, sizes). The metrics of interest will be: 1) the run-time of the program as the implementation of different allocation policies may result in different amounts of memory allocation overhead, and 2) fragmentation (i.e. the amount of unallocated data segment space divided by total data segment space). In order to capture #2, you should also implement two additional library functions:

```
unsigned long get_data_segment_size(); //in bytes
```

```
unsigned long get_data_segment_free_space_size(); //in byte
```

The functions must include the space used for metadata also.

```
get_data_segment_size() = entire heap memory (this includes memory used  
to save metadata)
```

```
get_data_segment_free_space_size() = size of the "free list" = (actual  
usable free space + space occupied by metadata) of the blocks in your  
free list
```

Then, using these functions, you can use the included test programs as well as test programs of your own creation to evaluate and compare the algorithms.

Hints

Hint: For implementing malloc(), you should become familiar with the sbrk() system call. This system call is useful for: 1) returning the address that represents the current end of the processes data segment (called program break), and 2) growing the size of the processes data segment by the amount specified by “increment”.

```
void *sbrk(intptr_t increment);
```

Hint: A common way to implement malloc() / free() and manage the memory space is to **keep a data structure that represents a list of free memory regions**(think about why not keep a list of all memory regions and the time complexity for your implementation). This collection of free memory ranges would change as malloc() and free() are called to allocate and release regions of memory in the process data segment. You may design and implement your malloc and free using structures and state tracking as you see best fit.

The Starter Kit

A starter kit is included in a file on Canvas in a folder under Files: **homework1-kit.tgz**

This archive can be downloaded to a local computer and transferred to the VM using `scp` for *nix/Mac systems(the easiest way) or retrieved on the command-line using `wget` if hosted on a website or by using a file transfer app on Windows (MobaXterm for Windows has built-in file transfer with drag and drop). It can then be extracted using `tar xvzf homework1-kit.tgz`.

The kit includes:

- **Makefile:** A sample Makefile for libmymalloc.so.
- **general_tests/:** General correctness test, see README.txt for details.
- **alloc_policy_tests/:** Allocation policy test cases, see README.txt for details.
- **alloc_policy_tests_osx/:** Same thing adapted for MacOS.
NOTE: Mac OS is not the grading platform; these files are provided as a convenience. Additionally, you may need to change `#include "time.h"` to `#include "sys/time.h"`

NOTE: These test files are provided for debugging. Passing these cases does not guarantee the correctness of their code. It is your responsibility to add more test cases to thoroughly test their code.

Testing your system

tl;dr: Write your own tests! At least a basic print data structure and minimal testing can help you catch several bugs.

Code is provided for minimal testing, but the provided materials will not exhaustively evaluate the correctness of your implementation. It is recommended to create your own test software that uses your library, both to aid during development and to ensure correctness in a variety of situations. If you'd like a general introduction to developing test cases, see [Software Testing by Sarah Heckman](#).

For debugging make sure to use tools such as gdb and valgrind. Please contact TAs and get help if you are not comfortable with using these tools. Also, please use the relevant online resources that may help you. Make sure to check for any errors in valgrind as these may be uncaught bugs in your program.

A refresher on gdb in emacs

- In your Makefile(s), change -O3 to -ggdb3. -O3 is an optimization flag and -ggdb3 is a debugging flag. -ggdb3 is required to let gdb add information to the compiled binary that will help it pinpoint exact line numbers, function names, etc. (Note: Please change flag back to -O3 in your submission.)
- Here is a link to how you may use gdb in emacs: [AoD_1: Debugging with GDB in Emacs](#) You may have a better reference video, please feel free to share that with others.
- In emacs open any file in the directory where you have the binary you need to run. e.g. if it is mymalloc_test from general_tests, navigate to that directory and open a file.
- Then run M-x gdb
- Then type (r) for run.
- Your code ends in a seg-fault--If you have followed all the above steps, gdb in emacs already shows you the line on which the seg-fault has occurred.
- You can then run (bt) for backtrace which will show you which series of function calls led you here, e.g. line number in main test file which resulted in this function call.
- Please consult this GDB cheat sheet [GDB QUICK REFERENCE GDB Version 4](#) for more.

Valgrind

Some flags that may be useful with valgrind:

```
valgrind -v --leak-check=full --track-origins=yes ./program
```

Note: When you test your `large_range_alloc` with valgrind, it may report an error: `brk segment overflow`, you can ignore it since there is a data segment size limitation setting by valgrind.

Getting Help

Please read Eric Steven Raymond's [How To Ask Questions The Smart Way](#). If you don't know

how to do anything mentioned in the document, make sure to look it up online first and then seek further clarification with more details in your question. The TAs are always happy to help. However, remember that you need to help them help you.

Grading

Rubrics

- code correctness(54'), 12 test cases total (you can assume the input for all the test cases are always valid)
- code check(16'), splitting and merging blocks required
- report(30'), including implementation description, performance result presentation and result analysis

This project will not be graded against the performance unless the performance is too bad. i.e.: Some points may be deducted if running a single test case takes 1 or 2 minutes.

Auto-grader Policy

When you submit your code, you won't be able to see the autograder grade immediately. However, we will release a pre-grade of the autograder two days before the due date (on Jan 25 at 11:59 pm). If you'd like to view your autograder result, please make sure to submit your code before this pre-release date.

Detailed Submission Instructions

Please submit the following to **Gradescope**:

1. A written report called **report.pdf** that includes:
 - An overview of how you implemented the allocation policies,
 - Results from your performance experiments,
 - An analysis of the results (e.g. why do you believe you observed the results that you did for different policies with different malloc/free patterns, do you have recommendations for which policy seems most effective, etc.).
2. All source code in a directory named **"my_malloc"**
 - There should be a header file name **"my_malloc.h"** with the function definitions for all *_malloc() and *_free() functions.
 - You may implement these functions in **"my_malloc.c"**. If you would like to use different C source files, please describe what those are in the report.
 - There should be a **"Makefile"** which contains at least two targets:
 - 1) "all" should build your code into a shared library named "libmymalloc.so"
 - 2) "clean" should remove all files except for the source code files. The provided Makefile may be used as-is, expanded upon, or replaced entirely. If you have not compiled code into a shared library before, you should be able to find plenty of

information online, or just talk to the instructor or TA for guidance!

With this “Makefile” infrastructure, the test programs will be able to: 1) #include “my_malloc.h” and 2) link against libmymalloc.so (-lmymalloc), and then have access to the new malloc functions. Just like that, you will have created your own version of the malloc routines in the C standard library!

Here is the submission format. Please submit your files in the following way, otherwise, the reviewers will not be able to compile your code, resulting in a score of 0. You can omit the “report.pdf” when you submit your codes to the Extra Credit (The next page will introduce it).

Name	Size	Progress	✕
my_malloc.c	5.6 KB	<div></div>	
my_malloc.h	0.8 KB	<div></div>	
Makefile	0.2 KB	<div></div>	
libmymalloc.so	16.9 KB	<div></div>	
report.pdf	0.1 MB	<div></div>	

Extra Credit

The extra credit is worth **10%** of the assignment grade.

For the extra credit, there is a separate **Project 1 Midpoint Check Assignment** on Gradescope, which will be due on **1/20 11:59 pm**. If you wish to get this extra credit, please complete the **First Fit implementation of malloc() (make sure your free() also works)**, and submit your source code before the above due date (written report is not required for extra credit part, and please follow the **Detailed Submission Instructions** for the source code submission). You will need to pass all auto-grader test cases in order to get the credit.

This Project 1 Midpoint Check Assignment will not affect your Project 1 Assignment grade which is due on 1/25 11:59 pm. The auto-grader for the Project 1 Midpoint Check Assignment will only test your **First Fit implementation** while the Project 1 Assignment will test both First Fit and Best Fit.

FAQ

1. Memory block size:

Q: Shall we follow the 8 bytes alignment when we implement the malloc? What is the smallest size of the memory block when we allocate memories, shall it be 1 byte or other sizes?

A: It's OK If you don't follow the 8 bytes alignment.

2. Project 1 test cases:

Q: I wonder if we can assume the user always uses our functions correctly. Do we need to deal with the cases such as double freeing or freeing something that is not at the start of the block returned by our malloc?

A: You can assume all the malloc and free functions calls are normal and valid.

3. Q: I wonder if we are allowed to allocate extra space for meta-data in heap for every allocation request.

A: Yes, you can.