

ECE_650 Project2: Thread-Safe Malloc and Free

2025/1/30-31

Sienna Zheng (NetID: sz318)

Environment:

- C
- Linux VM

Introduction:

- Based on codes of project1.
- Implement two different thread-safe versions of `malloc()` and `free()` using the **best fit allocation policy**.

Version 1: Use **lock-based synchronization** to prevent race conditions. Functions:

- `void *ts_malloc_lock(size_t size);`
- `void ts_free_lock(void *ptr);`

Version 2: Implement without using **locks or semaphores**, except for acquiring a lock **only** when calling `sbrk()`. Functions:

- `void *ts_malloc_nolock(size_t size);`
- `void ts_free_nolock(void *ptr);`

Thread-Safe Model

- **Locking Version** (`ts_malloc_lock` / `ts_free_lock`)
 - Uses a global mutex (`pthread_mutex_t`) to ensure only one thread accesses the memory allocator at a time.
 - Prevents race conditions but may introduce contention in multi-threaded environments.
- **Non-Locking Version** (`ts_malloc_nolock` / `ts_free_nolock`)
 - Utilizes **thread-local storage** (`__thread`) to maintain separate free lists for each thread.
 - Avoids locks except when using `sbrk()`, which is inherently **not thread-safe**.
 - Reduces contention but may lead to memory fragmentation since memory cannot be shared across threads.

Performance Comparison

1. Experimental Results

Metric	Non-Locking Version (ts_malloc_nolock)	Locking Version (ts_malloc_lock)
Execution Time	0.098750 seconds	0.108753 seconds
Data Segment Size	42,281,160 bytes	43,265,248 bytes

2. Observations and Analysis

1. Execution Time

- The **non-locking version** runs **~9.2% faster** than the locking version.
- This confirms that avoiding locks reduces contention and improves performance in multi-threaded scenarios.

2. Memory Usage

- The **locking version has slightly higher memory usage (~2.3% more)**, likely due to global synchronization and less frequent thread-local fragmentation.

3. Conclusion

- **Performance:** The **non-locking version** is faster due to reduced synchronization overhead.
- **Memory Efficiency:** The **locking version** uses slightly more memory but may exhibit better memory reuse in the long run.
- Trade-offs:
 - **Use non-locking** for better performance in high-thread-count applications.
 - **Use locking** if memory efficiency and reuse are a priority.

Appendix: Code

my_malloc.h:

```
1  #include <unistd.h>  // for sbrk()
2  #include <stdio.h>
3  #include <stdlib.h>
4  #include <pthread.h>
5
6  #define META_SIZE sizeof(block_meta)
7
8  typedef struct _block_meta {
9      size_t size;
10     struct _block_meta *next, *prev;
11 } block_meta;
12
13 /**
14  * Thread-safe by using lock around malloc and free
15  */
```

```

16 pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
17
18 // head and tail ptr of free_list
19 block_meta *head = NULL;
20 block_meta *tail = NULL;
21
22 // for test
23 void print_free_list() {
24     block_meta *curr = head;
25     printf("Free list: ");
26     while (curr) {
27         printf("[%p: size=%zu] -> ", curr, curr->size);
28         curr = curr->next;
29     }
30     printf("NULL\n");
31 }
32
33 // allocate the whole free block by deleting it from the free_list
34 void allocate_block(block_meta *block) {
35     // 1. only one block in free list
36     if (head == block && tail == block) {
37         head = NULL;
38         tail = NULL;
39     }
40     // 2. block is head
41     else if (head == block) {
42         head = block->next;
43         block->next->prev = NULL;
44     }
45     // 3. block is tail
46     else if (tail == block) {
47         tail = block->prev;
48         block->prev->next = NULL;
49     }
50     // 4. block is in the mid part
51     else {
52         block->next->prev = block->prev;
53         block->prev->next = block->next;
54     }
55     block->prev = NULL;
56     block->next = NULL;
57 }
58
59 // split a free block into an allocated block followed by a free block
60 void split_block(block_meta *block, size_t size) {
61     block_meta *new_block = (block_meta *)((char *)block + META_SIZE + size);
62     // char offset
63     new_block->size = block->size - size - META_SIZE;
64
65     // 1. only one block in the free list
66     if (head == block && tail == block) {
67         head = new_block;
68         tail = new_block;
69         new_block->prev = NULL;
70         new_block->next = NULL;

```

```

71 // 2. block is head
72 else if (head == block) {
73     if (block->next) block->next->prev = new_block;
74     new_block->next = block->next;
75     new_block->prev = NULL;
76     head = new_block;
77 }
78 // 3. block is tail
79 else if (tail == block) {
80     if (block->prev) block->prev->next = new_block;
81     new_block->prev = block->prev;
82     new_block->next = NULL;
83     tail = new_block;
84 }
85 // 4. block is in the mid part
86 else {
87     new_block->prev = block->prev;
88     new_block->next = block->next;
89     if (block->prev) block->prev->next = new_block;
90     if (block->next) block->next->prev = new_block;
91 }
92
93 // allocated block
94 block->size = size;
95 block->prev = NULL;
96 block->next = NULL;
97 }
98
99 // allocate space on heap. An allocated block always needs a meta
100 block_meta *request_space(size_t size) {
101     void *requested = sbrk(size + META_SIZE);
102     // allocation failure
103     if (requested == (void *) -1) {
104         fprintf(stderr, "Error: allocation failure");
105         exit(EXIT_FAILURE);
106     };
107
108     block_meta *block = (block_meta *)requested;
109     block->size = size;
110     block->next = NULL;
111     block->prev = NULL;
112
113     //data_segment_size += size + META_SIZE;
114
115     return block;
116 }
117
118 // only need to merge with direct prev and next free blocks
119 void merge_blocks(block_meta *block) {
120     if (!head || !block) return;
121
122     if (block->next && (char *)block->next == (char *)block + META_SIZE +
        block->size) {
123         block_meta * next_block = block->next;
124         block->size += next_block->size + META_SIZE;
125         // renew tail

```

```

126     if (next_block == tail) {
127         tail = block;
128     }
129     block->next = next_block->next;
130     // here block->next might be NULL
131     if (next_block->next) {
132         next_block->next->prev = block;
133     }
134     next_block->prev = NULL;
135     next_block->next = NULL;
136 }
137
138 if (block->prev && (char *)block == (char *) (block->prev) + META_SIZE +
block->prev->size) {
139     block_meta * prev_block = block->prev;
140     prev_block->size += block->size + META_SIZE;
141     if (block == tail) {
142         tail = prev_block;
143     }
144     prev_block->next = block->next;
145     if (block->next) block->next->prev = prev_block;
146     block->prev = NULL;
147     block->next = NULL;
148 }
149 }
150
151
152 void bf_free(void *ptr) {
153     if (!ptr) return;
154
155     block_meta *block = (block_meta*)ptr - 1;    // meta address
156     // 1. head = null, empty free list
157     if (head == NULL) {
158         head = block;
159         tail = block;
160         block->prev = NULL;
161         block->next = NULL;
162     }
163     // 2. freed block before head, block becomes the new head
164     else if (block < head) {
165         block->next = head;
166         head->prev = block;
167         block->prev = NULL;
168         head = block;
169     }
170     // 3. freed block after tail
171     else if (block > tail) {
172         block->prev = tail;
173         tail->next = block;
174         block->next = NULL;
175         tail = block;
176     }
177     // 4. somewhere between head and tail
178     else {
179         block_meta *curr = head;
180         // find two free blocks right before and right after block

```

```

181     block_meta *front = NULL, *back = NULL;
182     while (curr) {
183         if (curr < block) front = curr;
184         if (curr > block) {
185             back = curr;
186             break;
187         }
188         curr = curr->next;
189     }
190     if (front -> next != back) {
191         printf("front -> next != back\n");
192     }
193     block->prev = front;
194     block->next = back;
195     front->next = block;
196     back->prev = block;
197 }
198
199 merge_blocks(block);
200 }
201
202
203 // find a suitable free block according to bf strategy
204 block_meta * find_bf_block(size_t size) {
205     block_meta * curr = head;
206     block_meta * res = NULL;
207     while (curr) {
208         // a new meta is always needed for allocation
209         if (curr->size >= size) {
210             // find the smallest space bigger equal than size
211             if (res == NULL || curr->size < res->size) {
212                 res = curr;
213                 // already found the best one
214                 if (curr->size == size) break;
215             }
216         }
217         curr = curr->next;
218     }
219     return res;
220 }
221
222 // Best Fit malloc/free
223 void *bf_malloc(size_t size) {
224     if (size == 0) return NULL;
225
226     block_meta *block = find_bf_block(size);
227     if (block == NULL) { // no adequate free block, expand the heap
228         block = request_space(size);
229         // if (!block) return NULL; // heap expansion failure
230     } else { // adequate free block founded
231         // 1. no need to split
232         if (block->size <= size + META_SIZE) {
233             allocate_block(block);
234         }
235         // 2. split the block
236         else {

```

```

237     split_block(block, size);
238 }
239 }
240 return (void *) (block + 1); // return the actual data's address by
    skipping block_meta
241 // ptr addition: 1 refers to 1 ptr's size, aka 1 block_meta size.
242 }
243
244 //Thread Safe malloc/free: locking version
245 void *ts_malloc_lock(size_t size) {
246     pthread_mutex_lock(&lock);
247     void * ptr = bf_malloc(size);
248     pthread_mutex_unlock(&lock);
249     return ptr;
250 }
251
252 void ts_free_lock(void *ptr) {
253     pthread_mutex_lock(&lock);
254     bf_free(ptr);
255     pthread_mutex_unlock(&lock);
256 }
257
258 /**
259  * Thread-safe using no lock:
260  * __thread for free_list
261  * lock only for sbrk()
262  */
263 pthread_mutex_t sbrk_lock = PTHREAD_MUTEX_INITIALIZER;
264
265 // head and tail ptr of free_list
266 __thread block_meta *head_nolock = NULL;
267 __thread block_meta *tail_nolock = NULL;
268
269 // allocate the whole free block by deleting it from the free_list
270 void allocate_block_nolock(block_meta *block) {
271     // 1. only one block in free list
272     if (head_nolock == block && tail_nolock == block) {
273         head_nolock = NULL;
274         tail_nolock = NULL;
275     }
276     // 2. block is head
277     else if (head_nolock == block) {
278         head_nolock = block->next;
279         block->next->prev = NULL;
280     }
281     // 3. block is tail
282     else if (tail_nolock == block) {
283         tail_nolock = block->prev;
284         block->prev->next = NULL;
285     }
286     // 4. block is in the mid part
287     else {
288         block->next->prev = block->prev;
289         block->prev->next = block->next;
290     }
291     block->prev = NULL;

```

```

292     block->next = NULL;
293 }
294
295 // split a free block into an allocated block followed by a free block
296 void split_block_nolock(block_meta *block, size_t size) {
297     block_meta *new_block = (block_meta *)((char *)block + META_SIZE + size);
298     // char offset
299     new_block->size = block->size - size - META_SIZE;
300
301     // 1. only one block in the free list
302     if (head_nolock == block && tail_nolock == block) {
303         head_nolock = new_block;
304         tail_nolock = new_block;
305         new_block->prev = NULL;
306         new_block->next = NULL;
307     }
308     // 2. block is head
309     else if (head_nolock == block) {
310         if (block->next) block->next->prev = new_block;
311         new_block->next = block->next;
312         new_block->prev = NULL;
313         head_nolock = new_block;
314     }
315     // 3. block is tail
316     else if (tail_nolock == block) {
317         if (block->prev) block->prev->next = new_block;
318         new_block->prev = block->prev;
319         new_block->next = NULL;
320         tail_nolock = new_block;
321     }
322     // 4. block is in the mid part
323     else {
324         new_block->prev = block->prev;
325         new_block->next = block->next;
326         if (block->prev) block->prev->next = new_block;
327         if (block->next) block->next->prev = new_block;
328     }
329
330     // allocated block
331     block->size = size;
332     block->prev = NULL;
333     block->next = NULL;
334 }
335
336 // allocate space on heap. An allocated block always needs a meta
337 block_meta *request_space_nolock(size_t size) {
338     pthread_mutex_lock(&sbrk_lock);
339     void *requested = sbrk(size + META_SIZE);
340     pthread_mutex_unlock(&sbrk_lock);
341
342     // allocation failure
343     if (requested == (void *) -1) {
344         fprintf(stderr, "Error: allocation failure");
345         exit(EXIT_FAILURE);
346     }

```



```

347     block_meta *block = (block_meta *)requested;
348     block->size = size;
349     block->next = NULL;
350     block->prev = NULL;
351
352     //data_segment_size += size + META_SIZE;
353
354     return block;
355 }
356
357 // only need to merge with direct prev and next free blocks
358 void merge_blocks_nolock(block_meta *block) {
359     if (!head_nolock || !block) return;
360
361     if (block->next && (char *)block->next == (char *)block + META_SIZE +
block->size) {
362         block_meta * next_block = block->next;
363         block->size += next_block->size + META_SIZE;
364         // renew tail
365         if (next_block == tail_nolock) {
366             tail_nolock = block;
367         }
368         block->next = next_block->next;
369         // here block->next might be NULL
370         if (next_block->next) {
371             next_block->next->prev = block;
372         }
373         next_block->prev = NULL;
374         next_block->next = NULL;
375     }
376
377     if (block->prev && (char *)block == (char *) (block->prev) + META_SIZE +
block->prev->size) {
378         block_meta * prev_block = block->prev;
379         prev_block->size += block->size + META_SIZE;
380         if (block == tail_nolock) {
381             tail_nolock = prev_block;
382         }
383         prev_block->next = block->next;
384         if (block->next) block->next->prev = prev_block;
385         block->prev = NULL;
386         block->next = NULL;
387     }
388 }
389
390
391 void bf_free_nolock(void *ptr) {
392     if (!ptr) return;
393
394     block_meta *block = (block_meta*)ptr - 1;    // meta address
395     // 1. head = null, empty free list
396     if (head_nolock == NULL) {
397         head_nolock = block;
398         tail_nolock = block;
399         block->prev = NULL;
400         block->next = NULL;

```

```

401     }
402     // 2. freed block before head, block becomes the new head
403     else if (block < head_nolock) {
404         block->next = head_nolock;
405         head_nolock->prev = block;
406         block->prev = NULL;
407         head_nolock = block;
408     }
409     // 3. freed block after tail
410     else if (block > tail_nolock) {
411         block->prev = tail_nolock;
412         tail_nolock->next = block;
413         block->next = NULL;
414         tail_nolock = block;
415     }
416     // 4. somewhere between head and tail
417     else {
418         block_meta *curr = head_nolock;
419         // find two free blocks right before and right after block
420         block_meta *front = NULL, *back = NULL;
421         while (curr) {
422             if (curr < block) front = curr;
423             if (curr > block) {
424                 back = curr;
425                 break;
426             }
427             curr = curr->next;
428         }
429         if (front -> next != back) {
430             printf("front -> next != back\n");
431         }
432         block->prev = front;
433         block->next = back;
434         front->next = block;
435         back->prev = block;
436     }
437
438     merge_blocks_nolock(block);
439 }
440
441
442 // find a suitable free block according to bf strategy
443 block_meta * find_bf_block_nolock(size_t size) {
444     block_meta * curr = head_nolock;
445     block_meta * res = NULL;
446     while (curr) {
447         // a new meta is always needed for allocation
448         if (curr->size >= size) {
449             // find the smallest space bigger equal than size
450             if (res == NULL || curr->size < res->size) {
451                 res = curr;
452                 // already found the best one
453                 if (curr->size == size) break;
454             }
455         }
456         curr = curr->next;

```

```

457     }
458     return res;
459 }
460
461 // Best Fit malloc/free
462 void *bf_malloc_nolock(size_t size) {
463     if (size == 0) return NULL;
464
465     block_meta *block = find_bf_block_nolock(size);
466     if (block == NULL) { // no adequate free block, expand the heap
467         block = request_space_nolock(size);
468         // if (!block) return NULL; // heap expansion failure
469     } else { // adequate free block founded
470         // 1. no need to split
471         if (block->size <= size + META_SIZE) {
472             allocate_block_nolock(block);
473         }
474         // 2. split the block
475         else {
476             split_block_nolock(block, size);
477         }
478     }
479     return (void *) (block + 1); // return the actual data's address by
    skipping block_meta
480     // ptr addition: 1 refers to 1 ptr's size, aka 1 block_meta size.
481 }
482
483 //Thread Safe malloc/free: non-locking version
484 void *ts_malloc_nolock(size_t size) {
485     return bf_malloc_nolock(size);
486 }
487 void ts_free_nolock(void *ptr) {
488     bf_free_nolock(ptr);
489 }

```


